

Credit Card Fraud Detection

Credit card fraud is when someone uses another person's credit card or account information to make unauthorized purchases or access funds through cash advances. Credit card fraud doesn't just happen online; it happens in brick-and-mortar stores, too.

```
In [47]: import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns
```

```
In [48]: data = pd.read_csv("/creditcardfraud/creditcard.csv")  
data.head()
```

```
Out[48]:
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22	V23	V24	V25
0	0.00	-1.36	-0.07	2.54	1.38	-0.34	0.46	0.24	0.10	0.36	...	-0.02	0.28	-0.11	0.07	0.13
1	0.00	1.19	0.27	0.17	0.45	0.06	-0.08	-0.08	0.09	-0.26	...	-0.23	-0.64	0.10	-0.34	0.17
2	1.00	-1.36	-1.34	1.77	0.38	-0.50	1.80	0.79	0.25	-1.51	...	0.25	0.77	0.91	-0.69	-0.33
3	1.00	-0.97	-0.19	1.79	-0.86	-0.01	1.25	0.24	0.38	-1.39	...	-0.11	0.01	-0.19	-1.18	0.65
4	2.00	-1.16	0.88	1.55	0.40	-0.41	0.10	0.59	-0.27	0.82	...	-0.01	0.80	-0.14	0.14	-0.21

5 rows × 31 columns

Exploratory Data Analysis

```
In [49]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
 #   Column  Non-Null Count  Dtype  
---  -- 
 0   Time     284807 non-null   float64
 1   V1       284807 non-null   float64
 2   V2       284807 non-null   float64
 3   V3       284807 non-null   float64
 4   V4       284807 non-null   float64
 5   V5       284807 non-null   float64
 6   V6       284807 non-null   float64
 7   V7       284807 non-null   float64
 8   V8       284807 non-null   float64
 9   V9       284807 non-null   float64
 10  V10      284807 non-null   float64
 11  V11      284807 non-null   float64
 12  V12      284807 non-null   float64
 13  V13      284807 non-null   float64
 14  V14      284807 non-null   float64
 15  V15      284807 non-null   float64
 16  V16      284807 non-null   float64
 17  V17      284807 non-null   float64
 18  V18      284807 non-null   float64
 19  V19      284807 non-null   float64
 20  V20      284807 non-null   float64
 21  V21      284807 non-null   float64
 22  V22      284807 non-null   float64
 23  V23      284807 non-null   float64
 24  V24      284807 non-null   float64
 25  V25      284807 non-null   float64
 26  V26      284807 non-null   float64
 27  V27      284807 non-null   float64
 28  V28      284807 non-null   float64
 29  Amount    284807 non-null   float64
 30  Class     284807 non-null   int64 
dtypes: float64(30), int64(1)
memory usage: 67.4 MB
```

```
In [50]: pd.set_option("display.float", "{:.2f}".format)
data.describe()
```

Out[50]:

	Time	V1	V2	V3	V4	V5	V6	V7
count	284807.00	284807.00	284807.00	284807.00	284807.00	284807.00	284807.00	284807.00
mean	94813.86	0.00	0.00	-0.00	0.00	0.00	0.00	-0.00
std	47488.15	1.96	1.65	1.52	1.42	1.38	1.33	1.24
min	0.00	-56.41	-72.72	-48.33	-5.68	-113.74	-26.16	-43.56
25%	54201.50	-0.92	-0.60	-0.89	-0.85	-0.69	-0.77	-0.55
50%	84692.00	0.02	0.07	0.18	-0.02	-0.05	-0.27	0.04
75%	139320.50	1.32	0.80	1.03	0.74	0.61	0.40	0.57
max	172792.00	2.45	22.06	9.38	16.88	34.80	73.30	120.59

8 rows × 31 columns



In [51]: `data.isnull().sum().sum()`

Out[51]: 0

In [52]: `data.columns`

Out[52]: `Index(['Time', 'V1', 'V2', 'V3', 'V4', 'V5', 'V6', 'V7', 'V8', 'V9', 'V10', 'V11', 'V12', 'V13', 'V14', 'V15', 'V16', 'V17', 'V18', 'V19', 'V20', 'V21', 'V22', 'V23', 'V24', 'V25', 'V26', 'V27', 'V28', 'Amount', 'Class'], dtype='object')`

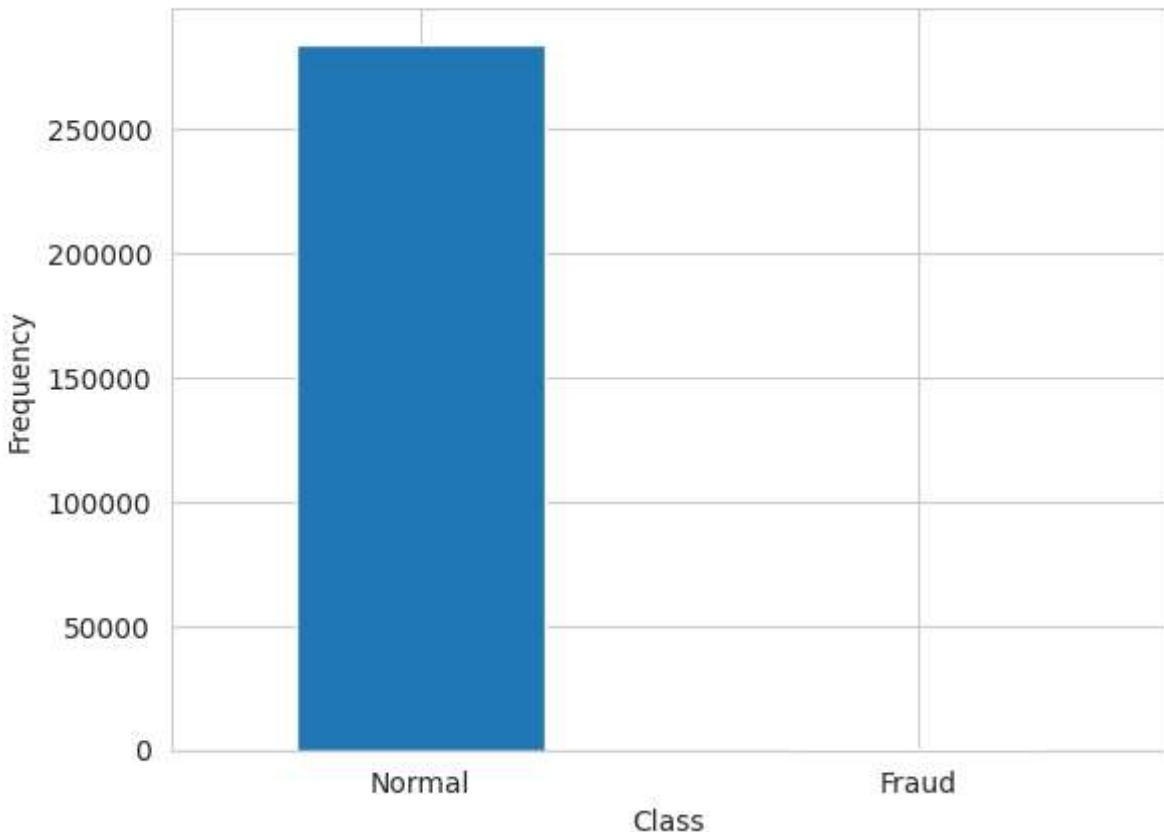
The only non-transformed variables to work with are:

- Time
- Amount
- Class (1: fraud, 0: not_fraud)

In [53]: `LABELS = ["Normal", "Fraud"]`

```
count_classes = pd.value_counts(data['Class'], sort = True)
count_classes.plot(kind = 'bar', rot=0)
plt.title("Transaction Class Distribution")
plt.xticks(range(2), LABELS)
plt.xlabel("Class")
plt.ylabel("Frequency");
```

Transaction Class Distribution



```
In [54]: data.Class.value_counts()
```

```
Out[54]: Class
0    284315
1      492
Name: count, dtype: int64
```

```
In [55]: fraud = data[data['Class']==1]
normal = data[data['Class']==0]

print(f"Shape of Fraudulent transactions: {fraud.shape}")
print(f"Shape of Non-Fraudulent transactions: {normal.shape}")
```

```
Shape of Fraudulent transactions: (492, 31)
Shape of Non-Fraudulent transactions: (284315, 31)
```

```
In [56]: pd.concat([fraud.Amount.describe(), normal.Amount.describe()], axis=1)
```

Out[56]:

	Amount	Amount
count	492.00	284315.00
mean	122.21	88.29
std	256.68	250.11
min	0.00	0.00
25%	1.00	5.65
50%	9.25	22.00
75%	105.89	77.05
max	2125.87	25691.16

In [57]: `pd.concat([fraud.Time.describe(), normal.Time.describe()], axis=1)`

Out[57]:

	Time	Time
count	492.00	284315.00
mean	80746.81	94838.20
std	47835.37	47484.02
min	406.00	0.00
25%	41241.50	54230.00
50%	75568.50	84711.00
75%	128483.00	139333.00
max	170348.00	172792.00

In [58]:

```
# plot the time feature
plt.figure(figsize=(14,10))

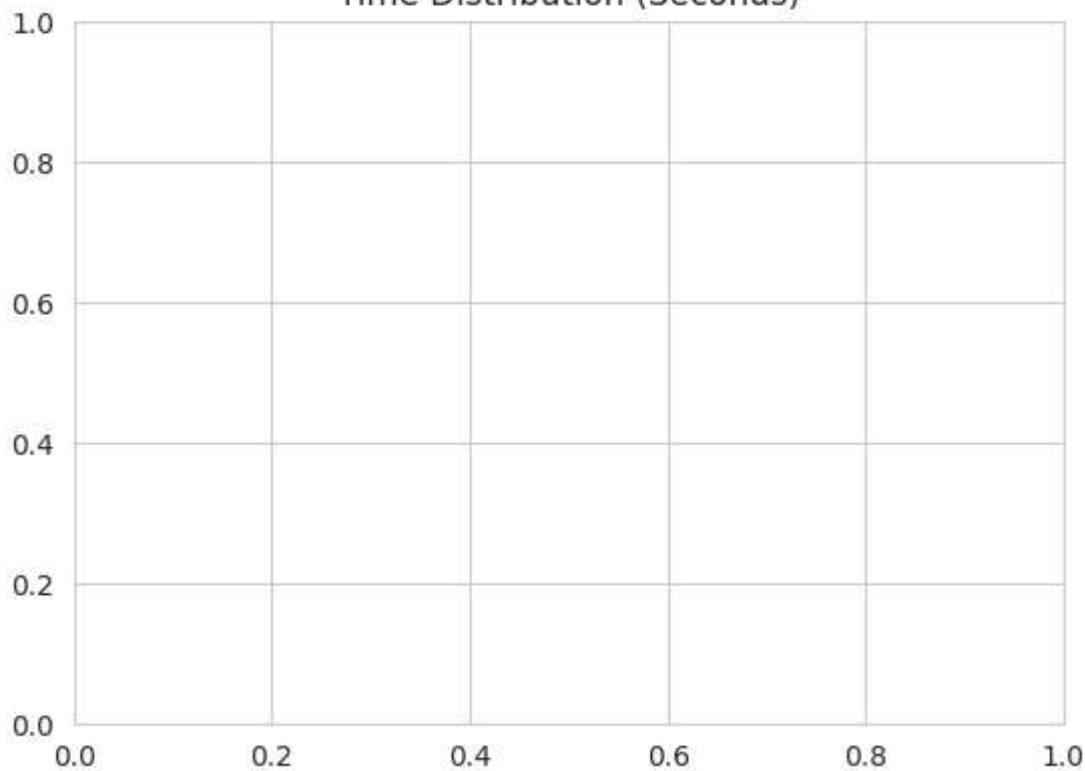
plt.subplot(2, 2, 1)
plt.title('Time Distribution (Seconds)')

sns.displot(data['Time'], color='blue');

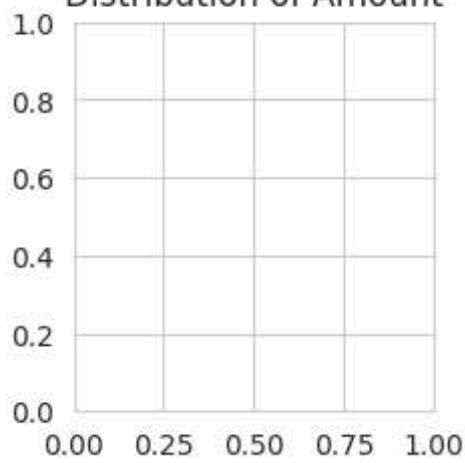
#plot the amount feature
plt.subplot(2, 2, 2)
plt.title('Distribution of Amount')
sns.displot(data['Amount'],color='blue');
```

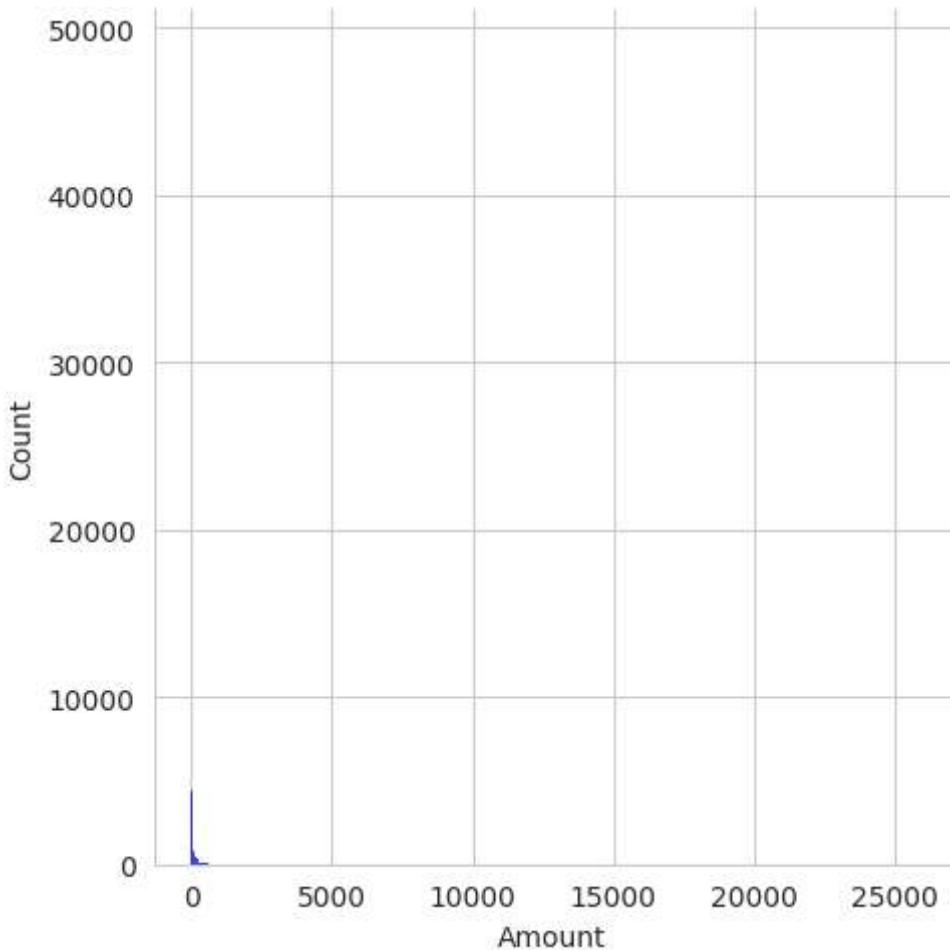
```
/opt/conda/lib/python3.10/site-packages/seaborn/axisgrid.py:118: UserWarning: The figure layout has changed to tight
    self._figure.tight_layout(*args, **kwargs)
/tmp/ipykernel_28/3161143837.py:10: MatplotlibDeprecationWarning: Auto-removal of overlapping axes is deprecated since 3.6 and will be removed two minor releases later; explicitly call ax.remove() as needed.
    plt.subplot(2, 2, 2)
/opt/conda/lib/python3.10/site-packages/seaborn/axisgrid.py:118: UserWarning: The figure layout has changed to tight
    self._figure.tight_layout(*args, **kwargs)
```

Time Distribution (Seconds)



Distribution of Amount



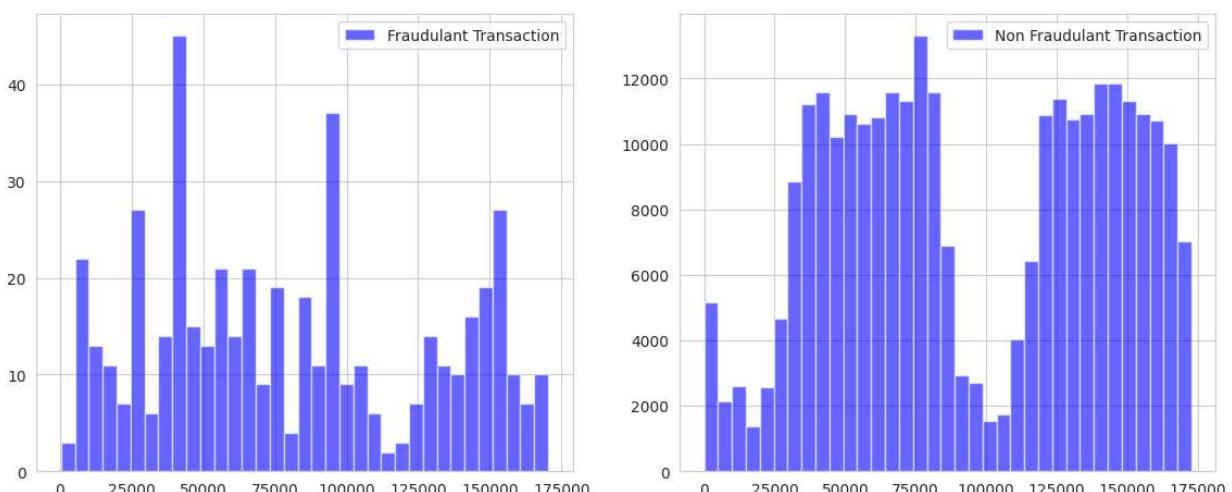


```
In [59]: # data[data.Class == 0].Time.hist(bins=35, color='blue', alpha=0.6)
plt.figure(figsize=(14, 12))

plt.subplot(2, 2, 1)
data[data.Class == 1].Time.hist(bins=35, color='blue', alpha=0.6, label="Fraudulent Transaction")
plt.legend()

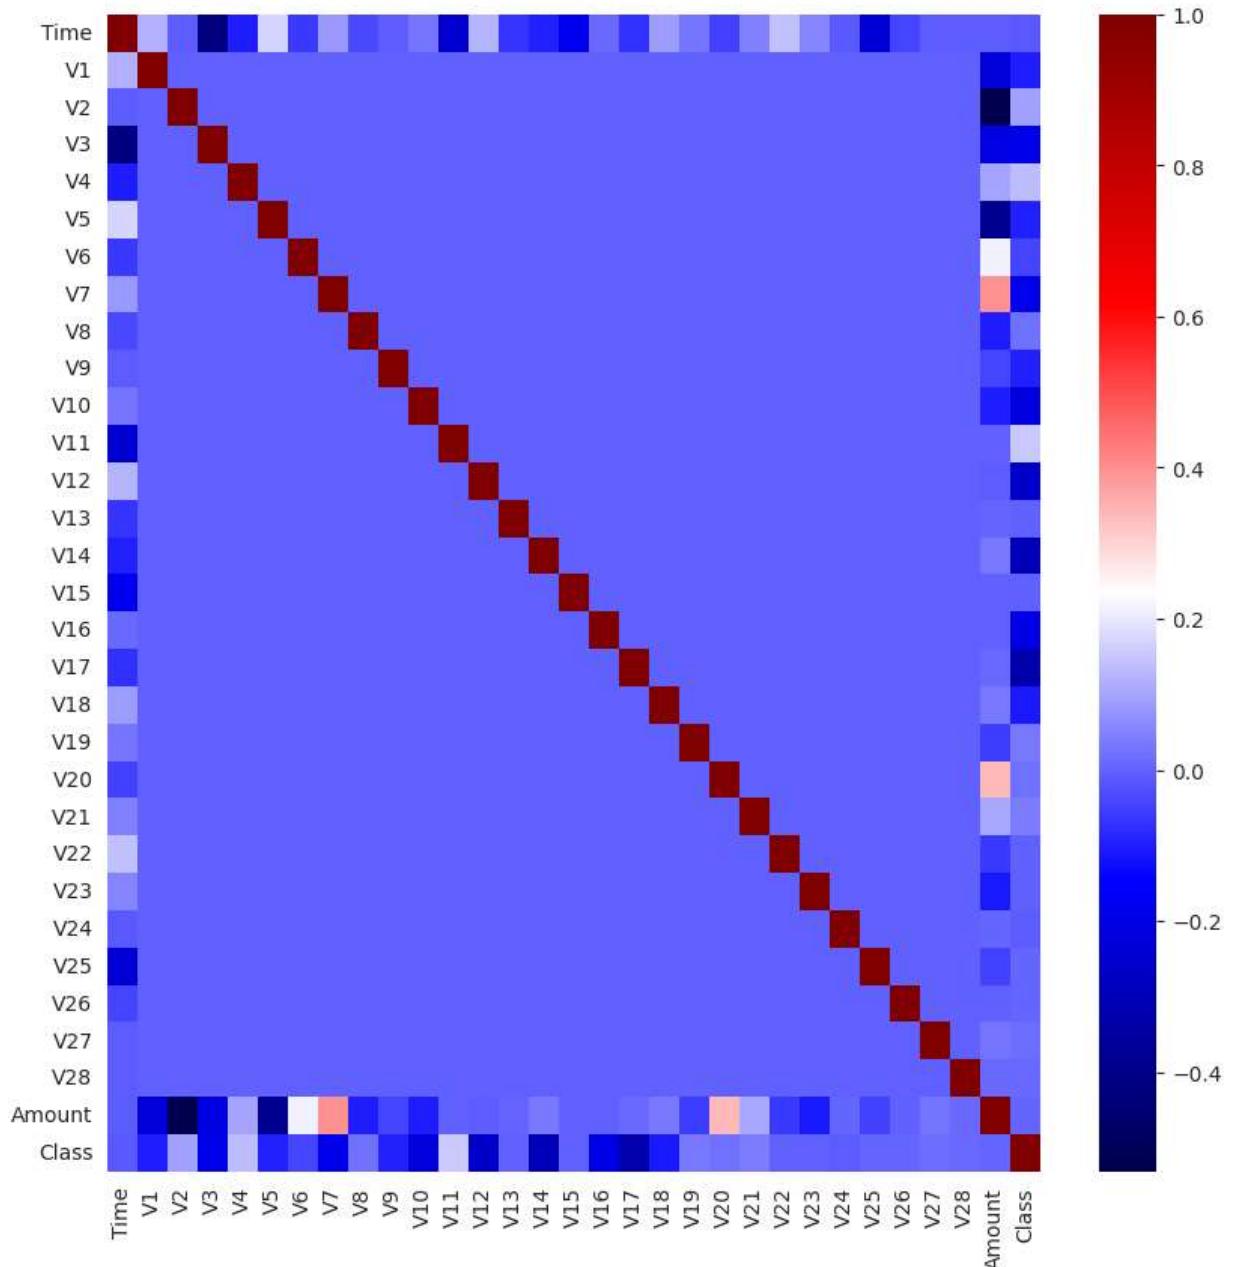
plt.subplot(2, 2, 2)
data[data.Class == 0].Time.hist(bins=35, color='blue', alpha=0.6, label="Non Fraudulent Transaction")
plt.legend()
```

Out[59]: <matplotlib.legend.Legend at 0x7c2f4330d540>



```
In [60]: # heatmap to find any high correlations
```

```
plt.figure(figsize=(10,10))
sns.heatmap(data=data.corr(), cmap="seismic")
plt.show();
```



Highest correlations come from:

- Time & V3 (-0.42)
- Amount & V2 (-0.53)
- Amount & V4 (0.4)

Data Pre-processing

Time and Amount should be scaled as the other columns.

```
In [61]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
```

```

scalar = StandardScaler()

X = data.drop('Class', axis=1)
y = data.Class

X_train_v, X_test, y_train_v, y_test = train_test_split(X, y,
                                                       test_size=0.3, random_state=42)
X_train, X_validate, y_train, y_validate = train_test_split(X_train_v, y_train_v,
                                                          test_size=0.2, random_state=42)

X_train = scalar.fit_transform(X_train)
X_validate = scalar.transform(X_validate)
X_test = scalar.transform(X_test)

w_p = y_train.value_counts()[0] / len(y_train)
w_n = y_train.value_counts()[1] / len(y_train)

print(f"Fraudulent transaction weight: {w_n}")
print(f"Non-Fraudulent transaction weight: {w_p}")

```

Fraudulent transaction weight: 0.0017994745785028623
 Non-Fraudulent transaction weight: 0.9982005254214972

```

In [62]: print(f"TRAINING: X_train: {X_train.shape}, y_train: {y_train.shape}\n{'*' * 55}")
print(f"VALIDATION: X_validate: {X_validate.shape}, y_validate: {y_validate.shape}\n{'*' * 55}")
print(f"TESTING: X_test: {X_test.shape}, y_test: {y_test.shape}")

TRAINING: X_train: (159491, 30), y_train: (159491,)

VALIDATION: X_validate: (39873, 30), y_validate: (39873,)

TESTING: X_test: (85443, 30), y_test: (85443,)

```

```

In [63]: from sklearn.metrics import accuracy_score, confusion_matrix, classification_report, f1_score

def print_score(label, prediction, train=True):
    if train:
        clf_report = pd.DataFrame(classification_report(label, prediction, output_dict=True))
        print("Train Result:\n====")
        print(f"Accuracy Score: {accuracy_score(label, prediction) * 100:.2f}%")
        print("____")
        print(f"Classification Report:\n{clf_report}")
        print("____")
        print(f"Confusion Matrix: \n {confusion_matrix(y_train, prediction)}\n")

    elif train==False:
        clf_report = pd.DataFrame(classification_report(label, prediction, output_dict=True))
        print("Test Result:\n====")
        print(f"Accuracy Score: {accuracy_score(label, prediction) * 100:.2f}%")
        print("____")
        print(f"Classification Report:\n{clf_report}")
        print("____")
        print(f"Confusion Matrix: \n {confusion_matrix(label, prediction)}\n")

```

Model Building

Artificial Neural Network (ANNs)

```
In [64]: from tensorflow import keras
```

```
model = keras.Sequential([
    keras.layers.Dense(256, activation='relu', input_shape=(X_train.shape[-1],)),
    keras.layers.BatchNormalization(),
    keras.layers.Dropout(0.3),
    keras.layers.Dense(256, activation='relu'),
    keras.layers.BatchNormalization(),
    keras.layers.Dropout(0.3),
    keras.layers.Dense(256, activation='relu'),
    keras.layers.BatchNormalization(),
    keras.layers.Dropout(0.3),
    keras.layers.Dense(1, activation='sigmoid'),
])

model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
<hr/>		
dense_4 (Dense)	(None, 256)	7936
batch_normalization_3 (BatchNormalization)	(None, 256)	1024
dropout_3 (Dropout)	(None, 256)	0
dense_5 (Dense)	(None, 256)	65792
batch_normalization_4 (BatchNormalization)	(None, 256)	1024
dropout_4 (Dropout)	(None, 256)	0
dense_6 (Dense)	(None, 256)	65792
batch_normalization_5 (BatchNormalization)	(None, 256)	1024
dropout_5 (Dropout)	(None, 256)	0
dense_7 (Dense)	(None, 1)	257
<hr/>		
Total params: 142,849		
Trainable params: 141,313		
Non-trainable params: 1,536		

```
In [ ]: METRICS = [
    keras.metrics.Accuracy(name='accuracy'),
    keras.metrics.FalseNegatives(name='fn'),
    keras.metrics.FalsePositives(name='fp'),
    keras.metrics.TrueNegatives(name='tn'),
    keras.metrics.TruePositives(name='tp'),
    keras.metrics.Precision(name='precision'),
    keras.metrics.Recall(name='recall')
]
```

```

model.compile(optimizer=keras.optimizers.Adam(1e-4), loss='binary_crossentropy', metrics=[f1])

callbacks = [keras.callbacks.ModelCheckpoint('fraud_model_at_epoch_{epoch}.h5')]
class_weight = {0:w_p, 1:w_n}

r = model.fit(
    X_train, y_train,
    validation_data=(X_validate, y_validate),
    batch_size=2048,
    epochs=300,
    #     class_weight=class_weight,
    callbacks=callbacks,
)

```

In [66]:

```
score = model.evaluate(X_test, y_test)
print(score)
```

```
2671/2671 [=====] - 8s 3ms/step - loss: 0.0035 - fn: 26.0000
- fp: 17.0000 - tn: 85290.0000 - tp: 110.0000 - precision: 0.8661 - recall: 0.8088
[0.0035336429718881845, 26.0, 17.0, 85290.0, 110.0, 0.8661417365074158, 0.80882352590
56091]
```

In [67]:

```
plt.figure(figsize=(12, 16))

plt.subplot(4, 2, 1)
plt.plot(r.history['loss'], label='Loss')
plt.plot(r.history['val_loss'], label='val_Loss')
plt.title('Loss Function evolution during training')
plt.legend()

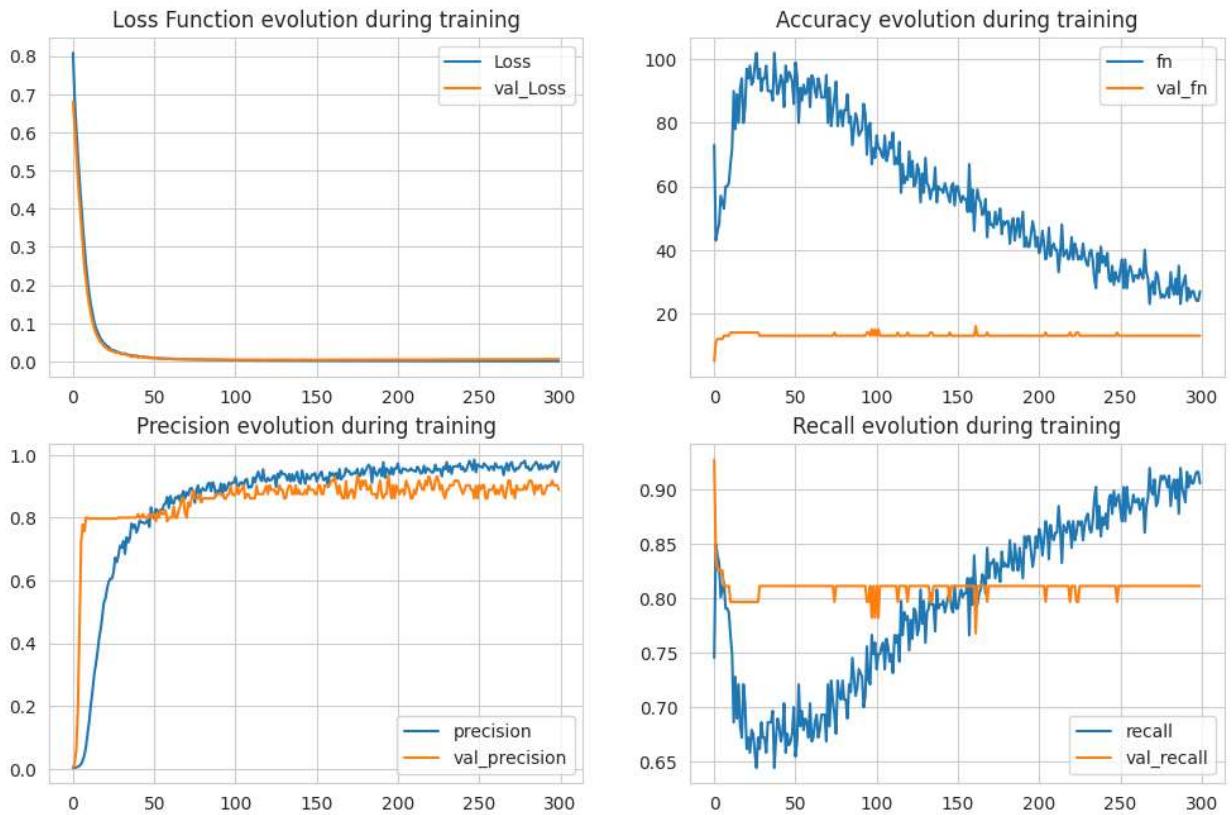
plt.subplot(4, 2, 2)
plt.plot(r.history['fn'], label='fn')
plt.plot(r.history['val_fn'], label='val_fn')
plt.title('Accuracy evolution during training')
plt.legend()

plt.subplot(4, 2, 3)
plt.plot(r.history['precision'], label='precision')
plt.plot(r.history['val_precision'], label='val_precision')
plt.title('Precision evolution during training')
plt.legend()

plt.subplot(4, 2, 4)
plt.plot(r.history['recall'], label='recall')
plt.plot(r.history['val_recall'], label='val_recall')
plt.title('Recall evolution during training')
plt.legend()
```

Out[67]:

```
<matplotlib.legend.Legend at 0x7c2f3d47d8d0>
```



```
In [68]: y_train_pred = model.predict(X_train)
y_test_pred = model.predict(X_test)

print_score(y_train, y_train_pred.round(), train=True)
print_score(y_test, y_test_pred.round(), train=False)

scores_dict = {
    'ANNs': {
        'Train': f1_score(y_train, y_train_pred.round()),
        'Test': f1_score(y_test, y_test_pred.round()),
    },
}
```

```
4985/4985 [=====] - 8s 2ms/step
2671/2671 [=====] - 4s 2ms/step
```

Train Result:

```
=====
```

Accuracy Score: 99.99%

Classification Report:

	0	1	accuracy	macro avg	weighted avg
precision	1.00	1.00	1.00	1.00	1.00
recall	1.00	0.96	1.00	0.98	1.00
f1-score	1.00	0.98	1.00	0.99	1.00
support	159204.00	287.00	1.00	159491.00	159491.00

Confusion Matrix:

```
[[159203 1]
 [ 11 276]]
```

Test Result:

```
=====
```

Accuracy Score: 99.95%

Classification Report:

	0	1	accuracy	macro avg	weighted avg
precision	1.00	0.87	1.00	0.93	1.00
recall	1.00	0.81	1.00	0.90	1.00
f1-score	1.00	0.84	1.00	0.92	1.00
support	85307.00	136.00	1.00	85443.00	85443.00

Confusion Matrix:

```
[[85290 17]
 [ 26 110]]
```

XGBoost

```
In [69]: from xgboost import XGBClassifier

xgb_clf = XGBClassifier()
xgb_clf.fit(X_train, y_train, eval_metric='aucpr')

y_train_pred = xgb_clf.predict(X_train)
y_test_pred = xgb_clf.predict(X_test)

print_score(y_train, y_train_pred, train=True)
print_score(y_test, y_test_pred, train=False)

scores_dict['XGBoost'] = {
    'Train': f1_score(y_train,y_train_pred),
    'Test': f1_score(y_test, y_test_pred),
}
```

```
/opt/conda/lib/python3.10/site-packages/xgboost/sklearn.py:835: UserWarning: `eval_me
tric` in `fit` method is deprecated for better compatibility with scikit-learn, use `

warnings.warn(
```

Train Result:

=====

Accuracy Score: 100.00%

Classification Report:

	0	1	accuracy	macro avg	weighted avg
precision	1.00	1.00	1.00	1.00	1.00
recall	1.00	1.00	1.00	1.00	1.00
f1-score	1.00	1.00	1.00	1.00	1.00
support	159204	287	1.00	159491.00	159491.00

Confusion Matrix:

```
[[159204    0]
 [    0   287]]
```

Test Result:

=====

Accuracy Score: 99.96%

Classification Report:

	0	1	accuracy	macro avg	weighted avg
precision	1.00	0.95	1.00	0.97	1.00
recall	1.00	0.82	1.00	0.91	1.00
f1-score	1.00	0.88	1.00	0.94	1.00
support	85307	136	1.00	85443.00	85443.00

Confusion Matrix:

```
[[85301    6]
 [  25  111]]
```

Random Forest

```
In [70]: from sklearn.ensemble import RandomForestClassifier

rf_clf = RandomForestClassifier(n_estimators=100, oob_score=False)
rf_clf.fit(X_train, y_train)

y_train_pred = rf_clf.predict(X_train)
y_test_pred = rf_clf.predict(X_test)

print_score(y_train, y_train_pred, train=True)
print_score(y_test, y_test_pred, train=False)

scores_dict['Random Forest'] = {
    'Train': f1_score(y_train,y_train_pred),
    'Test': f1_score(y_test, y_test_pred),
}
```

Train Result:

=====

Accuracy Score: 100.00%

Classification Report:

	0	1	accuracy	macro avg	weighted avg
precision	1.00	1.00	1.00	1.00	1.00
recall	1.00	1.00	1.00	1.00	1.00
f1-score	1.00	1.00	1.00	1.00	1.00
support	159204	287		1.00	159491.00
					159491.00

Confusion Matrix:

```
[[159204      0]
 [      0     287]]
```

Test Result:

=====

Accuracy Score: 99.96%

Classification Report:

	0	1	accuracy	macro avg	weighted avg
precision	1.00	0.90	1.00	0.95	1.00
recall	1.00	0.82	1.00	0.91	1.00
f1-score	1.00	0.86	1.00	0.93	1.00
support	85307	136		1.00	85443.00
					85443.00

Confusion Matrix:

```
[[85295    12]
 [  25   111]]
```

CatBoost

```
In [ ]: from catboost import CatBoostClassifier
```

```
cb_clf = CatBoostClassifier()
cb_clf.fit(X_train, y_train)
```

```
In [72]: y_train_pred = cb_clf.predict(X_train)
y_test_pred = cb_clf.predict(X_test)
```

```
print_score(y_train, y_train_pred, train=True)
```

```
print_score(y_test, y_test_pred, train=False)
```

```
scores_dict['CatBoost'] = {
    'Train': f1_score(y_train,y_train_pred),
    'Test': f1_score(y_test, y_test_pred),
}
```

Train Result:

=====

Accuracy Score: 100.00%

Classification Report:

	0	1	accuracy	macro avg	weighted avg
precision	1.00	1.00	1.00	1.00	1.00
recall	1.00	1.00	1.00	1.00	1.00
f1-score	1.00	1.00	1.00	1.00	1.00
support	159204	287	1.00	159491.00	159491.00

Confusion Matrix:

```
[[159204    0]
 [    1   286]]
```

Test Result:

=====

Accuracy Score: 99.96%

Classification Report:

	0	1	accuracy	macro avg	weighted avg
precision	1.00	0.93	1.00	0.97	1.00
recall	1.00	0.82	1.00	0.91	1.00
f1-score	1.00	0.87	1.00	0.94	1.00
support	85307	136	1.00	85443.00	85443.00

Confusion Matrix:

```
[[85299    8]
 [  25  111]]
```

LightGBM

```
In [73]: from lightgbm import LGBMClassifier

lgbm_clf = LGBMClassifier()
lgbm_clf.fit(X_train, y_train)

y_train_pred = lgbm_clf.predict(X_train)
y_test_pred = lgbm_clf.predict(X_test)

print_score(y_train, y_train_pred, train=True)
print_score(y_test, y_test_pred, train=False)

scores_dict['LightGBM'] = {
    'Train': f1_score(y_train,y_train_pred),
    'Test': f1_score(y_test, y_test_pred),
}
```

Train Result:

=====

Accuracy Score: 99.58%

Classification Report:

	0	1	accuracy	macro avg	weighted avg
precision	1.00	0.23	1.00	0.62	1.00
recall	1.00	0.59	1.00	0.79	1.00
f1-score	1.00	0.33	1.00	0.67	1.00
support	159204.00	287.00	1.00	159491.00	159491.00

Confusion Matrix:

```
[[158652    552]
 [   119    168]]
```

Test Result:

=====

Accuracy Score: 99.50%

Classification Report:

	0	1	accuracy	macro avg	weighted avg
precision	1.00	0.16	0.99	0.58	1.00
recall	1.00	0.53	0.99	0.76	0.99
f1-score	1.00	0.25	0.99	0.62	1.00
support	85307.00	136.00	0.99	85443.00	85443.00

Confusion Matrix:

```
[[84942    365]
 [   64    72]]
```

Model Comparison

```
In [74]: scores_df = pd.DataFrame(scores_dict)

scores_df.plot(kind='barh', figsize=(15, 8))
```

```
Out[74]: <Axes: >
```

